

A Language and Optimizing Compiler for Iterative Stencil Loops

Abstract

Parallel hardware is now widely available, yet writing and optimizing parallel programs using sequential languages is difficult and error prone. Therefore, it is useful to consider new language constructs that make it easier for programmers to express parallel applications without the need to attend to all the complexity of writing parallel algorithms. Here we present a new language in which users can easily express computations that involve iterative stencil calculations. Such computations form an important subset of parallel algorithms and are widely used in scientific, engineering, and other important application areas.

Specifically, we present a novel high-level language called Stencil Language (SL). Using SL, a programmer need only express the central kernel of their computation, not an entire parallel program. We have also built a translator for SL capable of translating the SL program into code for different parallel programming models. We provide an initial translation into optimized code for the C++ extensions of the CUDA programming model for NVIDIA GPUs. Our CUDA optimizer is an extension of a previously proposed analytical model for CUDA stencil applications. Our model improves over the previous model by automatically utilizing both static and dynamic information about the application and the specific GPU environment in which the application is running. This allows the optimized version of the application to be derived without the need for the application developer to specify any execution environment specific information. Finally, we provide results that show that the code generated from SL for CUDA is up to 2.3X faster than naïve CUDA code across six applications and two generations of NVIDIA hardware.

1. Introduction

Iterative stencil loops (ISLs) [6] are a class of loops that are commonly used in fields including numerical simulations and signal processing. Such loops compute a time varying set of values in an array. The value for each array element in each time step depends upon the value for that element and its nearby neighboring elements from the previous time step. The pattern of the accessed neighboring elements is called the *stencil*. This access pattern means that each element in the current time step can be calculated independently of all others. Therefore, techniques that take advantage of data parallelism are applicable to ISLs, and can often lead to dramatic speedups over serialized implementations. Yet, such data parallel implementations of ISLs present several challenges.

A typical implementation of an ISL on parallel hardware will divide the array into subsections called *tiles* and assign one or more tiles to each processor. A complication of this technique is that the neighboring elements near the edge of a tile assigned to one processor may reside in a tile assigned to another processor. Therefore, a certain amount of data sharing between the separate tiles assigned to different processors is required. Furthermore, on architectures where the cost of data communication between processors is expensive, an optimal ISL implementation will allow each processor to calculate several time steps for its tile(s) before communicating with other processors.

To accommodate this communication delay, the tiles assigned to different processors must now overlap. The amount of overlap

increases with the number of time steps between communication events. In addition, values of array elements in the overlap area must be redundantly calculated because each processor needs the boundary values locally for the intermediate time steps. On such hardware, the correct number of time steps to allow between processor synchronization is therefore a trade-off between the costs of the synchronization and associated data exchange, and the number of redundant values that are calculated.

These trade-offs are poorly understood and are highly architecture dependent. They are particularly challenging on new accelerator processors with novel memory hierarchies and synchronization support. GPUs are one prominent example. They have attracted great interest for general-purpose engineering and scientific computing due to their high data-parallel throughput and low cost.

What is needed is a way for users to express ISL computation in a way that allows them to focus on their application, and not on the parallel optimization techniques needed for fast performance. In this paper, we discuss our efforts to provide such a system. Our contributions are as follows:

- We have defined a *stencil language* called SL. This is a formal high-level language in which a programmer can describe an ISL computation. By design, a description written in the language must include all of the information necessary to calculate the correct results, and none of the information that would tie its performance to a particular architecture.
- We have created a front-end compiler for SL that can be extended with multiple back-ends. Based on a single source file, it could output efficient code for a variety of parallel architectures including CUDA [7, 13], OpenCL [17], OpenMP [2], or pthreads [11].
- We have provided an initial back-end for NVIDIA GPUs through source-to-source translation to their CUDA programming environment.
- Our CUDA code optimizer transparently calculates the optimal tile size and time steps between synchronization, and takes advantage of scratchpad memory. It extends the analytical model proposed by Meng and Skadron [9].
- Our system automatically collects all the necessary application specific and execution environment information needed by the model. Some is derived statically from the SL program, and the rest is derived dynamically at runtime.
- Dynamic acquisition of the execution environment information allows the generated code to calculate good values on a variety of GPU cards. We demonstrate this capability by using our optimizer developed on the GT200 series of GPUs on the newer Fermi architecture.

The remainder of this paper is organized as follows. We first provide a high level overview of our system in Section 2, describing the various components and their information flow. Then, in Section 3, we define the ISL terminology necessary for understanding the details of our SL language and compiler, which are then presented in Sections 4 and 5, respectively. Section 6 discusses the optimized CUDA code that is generated by our compiler, and Section 7 presents the performance of that generated code. Section 8 compares our approach to prior work, while Section 9 concludes and discusses our plans for future work.

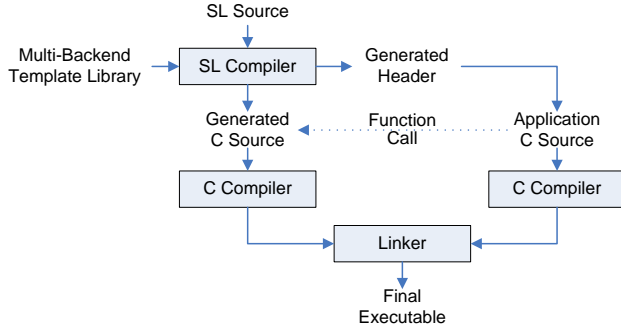


Figure 1. SL System Diagram showing information flow in a typical usage scenario. The SL source is combined by the SL compiler with the appropriate optimized back-end template. The resultant generated C subroutine for the ISL calculation is called by the rest of the user’s application code written in C or C++.

2. SL System Organization

Figure 1 shows the system organization of the SL compiler, and the information flow in a typical scenario. The application programmer specifies their ISL computation in an SL source file. The remainder of their application is specified in C or C++. In a simple scenario, the application code will set up the array upon which the ISL will operate, and call the code generated by the SL compiler to do the actual ISL computation. The results of the ISL computation can then be used by the application as necessary. The SL compiler generates a C Header File for inclusion in the application code to declare the signature of the ISL subroutine(s).

The SL compiler takes in two sources of information. One is the SL source file that specifies the user’s ISL computation. In addition, the SL system contains a template library that includes hand optimized implementations of the general structure of an ISL computation for various runtime environments. These templates are written by the SL system developers, not the SL user. Based on the user specified target environment, the SL compiler selects the appropriate library template and modifies it to implement the user’s particular stencil calculation. The SL compiler outputs the modified template as C source code. The SL generated C source and the user’s application C source are each compiled and linked together into a final executable.

While the simple case shows a single SL computation in an application, many scientific applications may contain more than one ISL computation in a pipeline. This usage scenario is supported by our architecture. Each ISL computation is specified in its own SL source file. The user’s application code can then call the generated subroutine for each ISL computation in sequence, passing the results of one calculation as input into the next.

While Figure 1 shows the general system organization, to date the template library only contains optimized templates for the CUDA environment. These templates include the analytical model for estimating the runtime for the generated CUDA kernel, plus the code to dynamically gather the runtime environment specific information needed by the model, and finally calculate the optimal tile size and time steps between synchronizations. The CUDA template causes the SL compiler to generate CUDA source code (a .cu file) which is then fed into NVIDIA’s NVCC compiler driver.

3. Iterative Stencil Loops: Terms & Background

Iterative stencil loops (ISLs) [6] are a class of loops commonly used in numerical simulations and signal processing. ISLs usually operate on matrices with one, two, or three dimensions. An ISL

calculation takes an input matrix m_{in} and produces an output matrix m_{out} . The calculation has the following properties:

1. m_{in} and m_{out} have the same number of dimensions and are the same size.
2. Each value in m_{out} is calculated independently of each other value in m_{out} .
3. There is a single function which is used to calculate each value in m_{out} regardless of the coordinates of the value.

For example, suppose m_{in} is a two-dimensional matrix, and the stencil calculation is defined as “ $m_{out}[x, y]$ = the average of the north, south, east, and west neighbors of $m_{in}[x, y]$ ”. Then this function will be calculated for every cell in m_{out} .

The *stencil* is the set of cells in the input matrix that is used in the stencil calculation, relative to the coordinates of the cell that is being calculated. In the above example, the stencil is a set of four cells: the north, south, east, and west neighbors.

The process of calculating a complete output matrix is called an *iteration* or a *time step*. At the end of an iteration, the output matrix m_{out} is used as the input matrix to the next iteration. This loop can continue for as long as necessary. Some iterative stencil loops are defined to run for a predetermined number of time steps, while others are defined to run until the values converge.

Stencil calculations are data parallel. Since each value in a time step is calculated independently of each other, the values can be calculated in any order or at the same time. If the matrix has n cells and the computer has k *processing elements* (PEs) available, then each processor only needs to calculate a block of size n/k . This partitioning of the matrices among multiple processors is called *tiling*, and the block calculated by a PE is called a *tile*.

When a PE calculates a given tile for many iterations, that PE can take advantage of both spatial locality and temporal locality.

- **Spatial locality.** As a general rule, stencil calculations use each input value multiple times. For example, say the stencil is defined as the cell’s left and right neighbors. When a PE calculates the cell value at $x = i$, it reads that value of the right neighbor at $x = i + 1$. Then, when it calculates the cell value at $x = i + 2$, it reads the left neighbor, which is the value at $x = i + 1$ again. Thus, a PE can optimize this computation by keeping input values in a local cache or scratchpad.
- **Temporal locality.** If a PE generates an output value during iteration t , then that PE can use that value as an input value during iteration $t + 1$.

Unfortunately, if each PE is given only those array values that fall within its tile size, cells along the boundary of the tile cannot find all of their input values in the local processor cache. These cells must obtain some of their input values from array values in adjacent tiles assigned to other PEs. In the example above, the cell in the bottom left corner of a tile will require values for both its south and west neighbors from neighboring tiles. The region of data surrounding a tile that is needed from adjacent tiles is called its *ghost zone* or GZ.

In addition, in order to take advantage of temporal locality, each PE may want to compute more than one time interval with the array values it has stored locally. However, at each successive time interval, more and more values at the periphery of the PEs local store will become stale. For example, consider the 2D tile computation shown in Figure 2(a). It shows the shrinking number of valid values calculated in a tile over four time intervals. In the first time interval i , with a stencil size of one, the PE will require a GZ of size one all around the periphery of the tile. However, in calculating the second time interval $i + 1$, the values in the original GZ are still from time i and are now out of date. Therefore, the PE is only able to calculate a smaller number of valid values, and




Figure 2. (a) For a single tile in a 2D matrix, the number of valid calculated values decreases with each additional time interval calculated between synchronization events. After 4 time steps, only the inner most rectangle of n values are valid. (b) Two adjacent tiles shown from the plane of the 2D array with time in the vertical dimension. In the first time step (diagram bottom) there is substantial overlap in the ghost zones in order to have the resultant valid values just abut after 4 time steps (diagram top).

the size of the GZ has expanded by the size of the stencil on all sides of the tile. In the diagram, this progression of stale values and expanding GZs will continue inward through time interval $t + 3$, after which only the inner most rectangle will have valid values.

Figure 2(b) shows the overlap in GZs between two adjacent tiles for the same computation over four time intervals with time now shown in the vertical dimension. Each PE must start with GZ of size four in order to have valid values computed for the resultant tiles cover the entire input array. Because the size of the area of valid values for a PE shrink with time in this way, the number of time intervals each PE calculates between synchronization events is called the *pyramid height* or PH of the ISL. In general, the size of the required GZ is equal to twice the size of the stencil times the PH. In addition, the values for cells in all but the initial GZ are redundantly calculated by adjacent PEs.

As a result of the required data exchange between PEs described above, ISL algorithms may spend considerable time stalled due to communication and synchronization delays. Because synchronization events can have high latencies, it is often better to reduce the number of synchronizations between PEs by increasing the PH and therefore the size of the GZ. However, as we have seen, larger GZs result in larger data exchanges at each synchronization, and also lead to redundant cell value calculations. Therefore, the choice of the optimal PH for an ISL is a trade-off between these two competing costs/benefits. Yet, the performance of ISLs depends critically upon the correct choice. Therefore, there have been several attempts to optimize these trade offs, such as the study by Meng and Skadron [9] in which they examined the problem of writing efficient ISL code in the CUDA language to run on an NVIDIA GPU. They defined a mathematical model for calculating the opti-

```
NumDimensions 2
StencilSize (1, 1)
DataType float
FunctionName runHotspot

ScalarVariables (
    float Rx, float Ry, float Rz
)
CellValue {
    float pvalue, value, term1, term2, term3, sum;
    pvalue = read(y * input_size.x + x);
    value = get(0, 0);
    term1 = (get(0, 1) + get(0, -1) - value) / Ry;
    term2 = (get(1, 0) + get(-1, 0) - value) / Rx;
    term3 = (80.0 - value) / Rz;
    sum = pvalue + term1 + term2 + term3;
    return(value + sum);
}
EdgeValue {
    return value;
}
```

Figure 3. The Hotspot.sl stencil language file. This example shows a 2D stencil calculation where each cell value is based on the values of that cell’s immediate neighbors. The read-only data is interpreted as a 2D matrix in row-major order which indicates that some areas of the chip inherently run hotter then other areas.

mal PH for an ISL targeted for that runtime environment. We have built upon the results of that study, and have fully automated that optimization process in our SL compiler, as we demonstrate in the following sections.

4. Stencil Language Definition

Figure 3 shows an SL description of the “Hotspot” calculation. This code calculates the temperature patterns in a two-dimensional microchip where the current state is a function of the previous state, the edge values, and a 2D read-only “power usage” matrix. We will examine the various aspects of this example in turn.

4.1 Syntax

A stencil language file is a text file that contains a set of key-value pairs. A key is an alphanumeric string. A value can be:

- an integer
- a name comprised of alphanumeric and underscore characters
- a list of names, inside parentheses, separated by commas
- a block of C code inside of curly braces

4.2 Content

The currently supported SL keywords are as follows.

- **NumDimensions** – A stencil may have 1, 2, or 3 dimensions.
- **StencilSize** – The size of the stencil in each dimension. If the stencil size is 1 in the x dimension, then the value of the cell at (x, y) may be based on the previous iteration’s values of the cells in the $x - 1$ and $x + 1$ coordinates.
- **DataType** – The data may be 32 bit or 64 bit integers or floating-point numbers. The values are `int`, `int64`, `float`, and `double`. Values may be signed or unsigned.
- **FunctionName** – The name of the function that will be exported from the generated source file. The first argument to the function, `data`, will be a matrix of the specified data type with the specified number of dimensions. As input to the function, `data` must contain the initial state of the problem. It must contain valid data in each cell. As output from the function, it will

contain the final data. The next n arguments to the function are the input sizes in each dimension. The next argument is the number of iterations to run.

- **ScalarVariables** – A list of numeric variables that will be arguments to the exported function after `data`, `x`, `y`, `z`, `iterations`. The caller will pass these arguments to the function, and then the stencil calculator will be able to use these variables (read-only) in its calculations.
- **CellValue** – A block of C code that will be run for every cell in the data set in every iteration. The code has access to the following variables and functions:
 - `x`, `y`, `z` – Coordinates of current cell in each dimension.
 - `iteration` – Iteration number. Iteration 0 is given as input, so this code will first be called with `iteration=1`.
 - `input_size` – a structure with `x`, `y`, and `z` fields, containing the total size of the input
 - All variables listed in the **ScalarVariables** field.
 - `get(...)` – a function that efficiently returns values from the stencil from the previous iteration. The parameters are relative to the current cell. For example, in a 2D stencil, the west and east neighbors can be retrieved with `get(-1, 0)` and `get(1, 0)`.
 - `read(...)` – a function to access read-only data (see below).

The block of code returns the new cell value.

- **EdgeValue** – A block of C code that returns the cell value of cells that are outside the bounds of the input. This code has access to the same variables as the **CellValue** code. Note that at least one of `x`, `y`, or `z` will be out of bounds – it will be either less than zero or greater than or equal to the size of that dimension. This code may not use the `get()` function. Instead, it can access the `value` variable, which will contain the value of the nearest cell in bounds from the previous iteration.

4.3 Read Only Data

The generated source file will export two functions. One of the functions will have the given function name. The other function will have the function name appended with `SetData`. For example, if the stencil language file says:

```
FunctionName runStencil
```

then the generated source file will export these two functions:

- `runStencil()`
- `runStencilSetData()`

The `SetData` function takes two arguments:

1. An array of elements of the specified data type
2. The number of elements in the array

If a program calls `runStencilSetData()` with an array, and then calls `runStencil()`, then the stencil calculator will have read-only access to this data in the array. The **CellValue** code can access the i th read-only data element by calling `read(i)`.

5. Stencil Language Compiler

We implemented the stencil language compiler in fewer than 600 lines of Java code. The compiler operates as follows:

1. Converts a stencil language file into a token stream
2. Parses the token stream into a symbol table
3. Reads a plain text template file

4. Outputs a new file which contains the template integrated with the stencil description

Step 1 is implemented by `Tokenizer.getToken()`. A token is either a string of alphanumeric characters or any symbol token that is in the language such as parentheses, commas, etc. Also, the tokenizer can recognize comments and blocks of C code.

Step 2 is implemented by `Stencil.parse()`. It runs a loop that reads a key and then parses a value, which may be an integer, a string, a list, or a block. The parser verifies that the key is defined by the language and that the value has the right type for that key, and then it adds the key value pair to a symbol table.

Step 3 is implemented by the `creator` method of the `Template` class. It stores the template text in the `Template` instance.

Step 4 is implemented by `applySymbolTable()` method. It looks for instances of `@name@` in the text. For each instance that it finds, it looks up `name` in the symbol table, and it replaces the variable with the value.

This design can easily support multiple architectures and optimizations. We have written templates for CUDA code for 1D, 2D, and 3D stencil calculations. These templates are written such that the **CellValue** code can be dropped-in unmodified. The template code is optimized for shared memory usage, and it calculates the optimal PH. We have also written a template that does not use shared memory or GZs to compare its performance as a naïve CUDA implementation to that of the optimized templates.

We are confident that an OpenMP template could be written without changes to either the stencil language or the compiler.

6. Generated CUDA Code

The following subsections describe the CUDA architecture, and the details of the SL CUDA template and optimizer.

6.1 CUDA Architecture Overview

The CUDA programming model [7, 13] is a set of extensions to C++ that NVIDIA Corporation created to allow their graphics cards (GPUs) to be used for general-purpose parallel computing. The overall architecture of a NVIDIA GPU [12] is shown in Figure 4. The model involves using the GPU as a co-processor for the CPU. A control thread on the CPU does some of the work for the application, but can call on the GPU to perform certain parallel calculations when desired. The parcel of work for the GPU is called a *kernel* and usually represents an inner loop that exhibits large degrees of data parallelism. The code for the kernel is written for one thread as if it were the only thread running, except that it is given a thread-id from which it can compute its position in the collection of threads launched in the kernel invocation. While the kernel code can be written rather simply for simple applications, the optimal use of the GPUs resources requires much more effort and complexity.

One source of the complexity for optimizing kernels comes from the specialized memory hierarchy on the device due to its primary use as a graphics accelerator. The hardware of the GPU is organized into a collection of PEs each of which is a SIMD core with 8 scalar processing elements. Both the organization of the GPU into separate PEs, and the related variety of different types of memory, are made visible to the kernel writer.

When calling a kernel, the user must specify the organization of the resultant threads into thread blocks, including the 3D shape of the thread block. GT200 series and earlier GPU cards support up to 512 threads per block. A thread block will be allocated by the CUDA runtime system to a particular PE on which it will execute. To minimize SIMD overhead, these threads are collected into smaller groups called *warps*. Each PE has a modest amount of very fast programmer controlled scratchpad memory that NVIDIA calls *shared memory* that can be shared by threads in a thread block,




Figure 4. NVIDIA GPU Diagram. The CPU control thread launches a kernel with many threads per thread block. Each thread block is assigned to a Multiprocessor (PE) containing 8 scalar processing elements and a local shared memory. The Device Memory is global to all thread blocks and in pre-Fermi GPUs is uncached.

but cannot be accessed by any other thread blocks. For the purposes of this paper, the only other memory structure available to threads running on the GPU is a large global memory accessible by all threads on the device. However, on pre-Fermi GPUs, this large memory is not cached, and has access latencies of 400 to 600 clock cycles. Therefore, a critical optimization for kernel writers is to properly group threads into blocks that can use the shared memory effectively.

In addition, the CUDA model does not support any global synchronization primitive between thread blocks during a kernel execution. Instead, global synchronization events require the completion of all running thread blocks and a return of control to the CPU control thread. A further complication is that no data stored in the shared memory of a PE is persistent across kernel calls. Such global synchronization events are therefore very costly.

The implication of this model for stencil applications is that the kernel author must be intimately familiar with the details of the CUDA runtime system in order to organize work units into thread blocks in such a way that usage of shared memory is maximized, and accesses to global memory are minimized. Key decisions in this allocation are precisely the tile size and PH of the stencil calculation that the optimization model is designed to make. In order to accurately make such decisions, it must be able to predict the latencies of the various operations involved in the kernel's calculations, and memory accesses on the GPU device upon which the application will run.

6.2 SL CUDA Template

In this section, we describe how our CUDA template code works. This template code is hand-written by the implementors of the SL system, and does not contain any application specific code. The application specific code from the SL specification is then automatically merged into the template during the SL compilation process to generate the final ISL code. The CUDA execution model involves both a control thread on the CPU and kernel code that is run by many threads on the GPU. Therefore, the CUDA template code includes both the CPU and GPU code needed to support the application specific code provided by SL.

The primary function in the template is a C function for the CPU control thread. First, it calculates the tile size and the PH. Then, it copies the input matrix from system memory into GPU memory. If an array of read-only data is used, it also copies that data into GPU memory. Finally, it runs a loop in which it invokes the CUDA kernel to calculate the output matrix from the input matrix. It runs this loop for the specified number of iterations, divided by the PH.

Each tile is processed by one CUDA thread block, and each cell in the tile is owned by a single thread. The input tile size will include GZ cells. The output tile size will be smaller. The actual tile size used is determined at application runtime using dynamically gathered information about the execution environment. For concreteness, consider a pre-Fermi CUDA card that supports up to 512 threads per thread block. In this case, the 3D input tile will be $8 \times 8 \times 8$, the 2D tile 22×22 , and the 1D tile 512 cells long.

The generated CUDA kernel includes both template code and application specific code. It operates as follows. The thread calculates the (x, y, z) coordinates of the cell in the matrix that it owns. It copies that cell's value from the input matrix into the thread block's shared memory. After each value has been copied into shared memory, the threads that own cells in the outermost GZ sleep. Those threads have completed their work. Each other thread calculates the new value for its cell. For each cell in its stencil, the thread gets the cell value from shared memory.

If the PH is greater than one, then the threads continue running. Each thread copies the value that it just computed into shared memory. Then, the threads in the second outermost GZ sleep. Their local copies of the neighboring cell values in the outermost GZ are now out-of-date as their values were calculated in an adjacent tile (see Figure 2). The remaining threads calculate new values. This process continues for the number of iterations specified by PH.

When this process is complete, if a thread owned a cell that was not in the GZ, then that cell copies the final value into device memory. If a thread owned a cell that was in the GZ, then that thread does not need to do anything, since we are guaranteed that some thread in some other thread block owned that same cell as part of its output tile, and that other thread will own the responsibility to write the value into device memory.

Inside of the CUDA kernel, the program calls the user-specified CellValue code to calculate new cell values, and it calls the user-specified EdgeValue code to get stencil values for cells along the edges of the matrix. All of the remainder of the kernel code is included in the template in the SL template library.

6.3 Pyramid Height Calculation

As mentioned earlier, the performance of stencil applications on parallel hardware is critically dependent on the proper relationship between tile size and GZ size or PH. To calculate the optimal PH, we extended and automated the CUDA analytical model written by Meng and Skadron. They developed a complex analytical model in MatLab that uses information from several sources to calculate the expected program runtime (in GPU clock cycles). Their model requires several pieces of information as input.

First, it needs information about the stencil application itself. Some information is static such as the data dimensionality, the stencil size (GZ), the number of global memory accesses required during movement of tile data into shared memory, and the number of global accesses for read-only data required during each cell calculation. Other information is dynamic such as the data set size and the number of instructions that each thread must execute (1) in the shared memory setup portion of the kernel, and (2) in the kernel cell calculation. Since they proposed but did not implement any code annotations for the stencil applications, many of the application specific parameters were hand coded in MatLab for the four particular test case applications. In addition, the instruction count

was gathered using the CUDA profiler on hand-coded implementations of the test applications, and fed to the model by hand.

Second, the specific properties of the target GPU are important. This includes the number of PEs, the number of concurrent blocks that can run on a PE, the number of threads allowed per block (tile), the time required for the GPU to perform a global synchronization, and several GPU memory performance parameters. These values were acquired from a number of means including hand coding some parameters based on the published characteristics of a few specific models of GPU, and some by data regression (curve fitting) of runtimes from specific applications.

We have modified and/or improved this model in several important ways. First, we ported the model from MatLab to C++ to include it in the application and perform the optimization at runtime. This allows gathering runtime information about the program in the current execution environment allowing it to be run on different GPUs without recompilation. The resulting code is better optimized for the specific characteristics of the particular hardware, not just the architectural features of the hardware. While this means that the time to perform the optimizations is included in the total runtime for the program, the difference in runtimes for the wrong optimizations will generally dwarf any cost of calculating the optimizations themselves. In addition, stencil applications typically take many iterations to converge, allowing more runtime over which to amortize the optimization costs.

Second, the static and architecture-independent application specific information about the dimensions of the data space, the stencil size, etc. now come directly from the SL program specification. These are compiled into the generated code for the ISL calculation.

Third, the remaining application-specific information is acquired dynamically at runtime. This includes the data set size, and the parameters that previously came from profiling such as the kernel thread setup time and cell value calculation time. The latter two numbers are acquired by running the application with PH of 1 and 2. From these numbers, the setup and cell calculation times are easily computed. Currently we throw away the application computation performed during these data gathering runs, but in the future, the time spent gathering this information could result in the completion of the first three iterations of the computation.

Fourth, the GPU device information is gathered through calls to various functions provided in the CUDA runtime libraries. In particular, we can acquire the number of PEs, the concurrent blocks per PE, and the maximum number of threads per block. We can also query the CUDA runtime system for the number of threads per block for the actual kernel of the running application. This can sometimes result in tighter restrictions on the allowed threads per block if, for example, the kernel requires more resources from the device than it can provide at the architecture maximum number of threads per block. This is particularly true for CUDA kernels that use a lot of registers. Knowing this number is critical to the optimization because on the CUDA architecture, stencil applications generally run faster with a larger tile size as long as the data set size is large enough to saturate the number of available PEs. Since the thread per block count determines the tile size, knowing this number precisely allows the optimizer to pick the largest tile size that the current GPU can handle. Future GPUs may allow the number of threads per block to increase beyond 512. For example, Fermi GPUs now allow 1024 threads per block. Gathering this information at run time allows generated code to calculate the optimal tile size for such cards without recompilation.

Fifth, the time required for a global synchronization is acquired by running a null kernel on the device with the same data set size parameters as are required to run the actual kernel.

Finally, Meng and Skadron originally did a manual gradient descent on runtimes to determine the best PH and tile size. Since

we can determine the best tile size as described earlier, we do an automatic gradient descent starting at PH of 1, and continue calculating the model latency until it is higher than the previous one, then use that PH for the remainder of the application iterations.

7. Experimental Results

Our primary emphases are the definition of a stencil language, providing a compiler for it, and optimizing the generated programs in terms of GPU shared memory, tile size, and PH with an extension of the analytical model of Meng and Skadron. Therefore, we first present the results for NVIDIA's GT200 series cards, which is the architecture for which their performance model was developed. The recent release of NVIDIA's Fermi series cards presented an opportunity to test the robustness to architectural changes of our model enhancement which utilizes dynamically-collected information about the execution environment. Therefore, we also tested on a GTX-480 card after making one minor change to our optimizer to better weight the global synchronization latency in the model. All SL optimized results shown below for both the GTX-280 and GTX-480 are from this updated version of the optimizer.

All runtimes include the CPU and GPU latencies for the ISL loop only. The time required to create and initialize arrays, and transfer arrays between main memory and the GPU global memory are not included to allow us to focus on the factors affecting the ISL performance. In addition, the reported runtimes for the SL optimized ISLs do not include the overhead needed to run the optimizer. Instead, the times are for the entire ISL loop running with the parameters determined by the optimizer. As stated above, the costs for running the optimizer are fixed, so the overhead for the optimizer as a percentage of the overall runtime can be arbitrarily large or small depending on the number of iterations in the actual ISL loop.

We use the arithmetic mean when aggregating results, and we use a SERPOP analysis strategy (see [8] for details).

7.1 GT200 Results

Our initial testing environment has 2 Intel Core2 Extreme X9770 processors running at 3.2GHz with 6MB of L2 cache and 4GB of DRAM running Linux 2.6.24. The graphics card is an NVIDIA GT200 series card - the GTX-280 - running graphics driver v3.1. GT200 cards have 30 PEs, each with 8 SIMD cores that handle up to 8 concurrent running blocks, with a 512 thread maximum block size. The GTX-280 runs at 1.3GHz, with 1GB of global memory. All code was compiled with the NVIDIA NVCC compiler driver which calls GNU tools (version 4.2.4) for compiling C++ code and linking. All code was compiled with -O3.

For our first test cases, we use the same four stencil applications as Meng and Skadron did in their original study. They are described below. All of the graphs of runtimes vs. PH have the same basic curve. Very low PHs are not large enough to take advantage of all available spatial and temporal locality, and perform worse than larger PHs, especially given the relatively large cost of global synchronizations in the CUDA architecture. However, beyond some optimal PH, the runtimes again increase due to the overhead of all the redundant local GZ value calculations that are thrown away. Another way to see this trade-off is that with increased PH, the number of threads per tile that calculate final cell values drops off at twice the size of the stencil per dimension per step in PH (see Figure 2). Therefore, for larger PHs the effective tile size gets smaller, and the total number of tiles, and therefore also the total number of threads, increases. Eventually, in the limiting case, the PH is so high that *no* useful cell values are calculated. Near this extreme, run times increase exponentially. For these reasons, all of the graphs of runtimes vs. PH will be bowl shaped. To better highlight the im-

Pathfinder Runtimes for 4 data set size

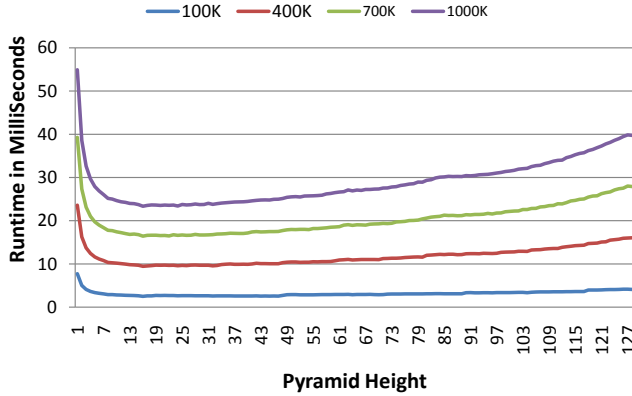


Figure 5. Runtimes for the Pathfinder 1D application vs. PH. The lowest portion of the curves are very flat. Therefore, even though our model did not pick the best PH for any data set size, the performance was within 0.99% and 2.79% of optimal.

portant portions of the curves near the optimal PH, in the following graphs we have not shown the extreme values at very high PH.

7.2 “Pathfinder”

The first test case is a 1D stencil application called “Pathfinder” that performs a dynamic programming calculation of the lowest cost of any vertical path through a 2D grid. The 2D input data set is read-only data in the CellValue computation. Each loop iteration calculates the minimum costs down another row in the 2D grid. The tiles are horizontal slices across the columns of the grid. A valid path is one which varies by at most one cell in each step from row to row. Therefore, the stencil size is 1. Figure 5 shows the runtimes vs. PH for Pathfinder for row widths of 100K, 400K, 700K, and 1M. As with all the applications, these large sizes are chosen to ensure that the GPU is saturated with threads. As this is a 1D application, the possible PHs vary from 1 to 255 on GT200 cards.

Notice that the valley for the main portion of this curve is very shallow. Therefore, it is hard for the model to accurately predict the exact best PH. And in fact, our model does not. The measured best PH is 16 for all data set sizes. However, our optimizer predicts 41, 21, 21, and 20 for the 100K, 400K, 700K and 1M data set sizes respectively. For the 400K data set size, this results in a performance loss of 5.8%, while for the others the performance loss is just under 1%. From these results, one can see that the shape of the interesting part of the curve is so flat that precise PH prediction is not necessary to achieve very positive results. For comparison, the poor choice of a PH of 1 for the 400K column data set would result in a performance loss of nearly 2.5X.

7.3 “Hotspot”

The second test case is a 2D stencil application called “Hotspot”. The stencil language description for this application is shown in Figure 3. The application calculates an ordinary differential equation that models heat dissipation in a conductive material. In this case, the modeled material is the silicon substrate of a computer chip [16]. On each time step of the computation, the heat generated by the chip components are injected through the modeled substrate. This is constant read-only data read for each cell on each iteration. As this is a 2D application, the theoretical maximum square tile size on a GT200 card is 22x22. This implies a possible PHs range from 1 to 10. The resultant runtimes vs. PH are shown in Figure 6.

HotSpot Runtimes for 4 data set sizes

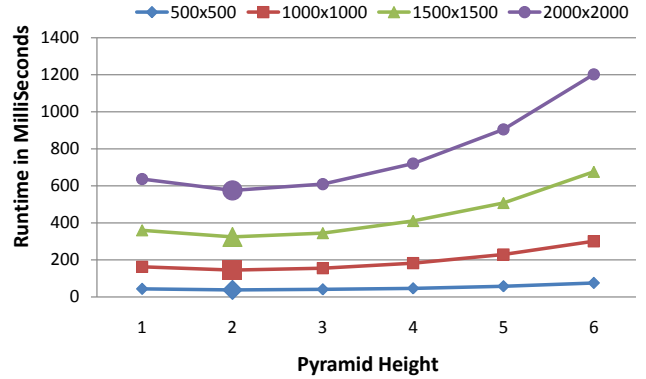


Figure 6. Runtimes for the Hotspot 2D stencil application vs. PH. As can be seen, the curve is bowl shape with a noticeable low point at PH 2. The effect is more pronounced with greater data set sizes. Our optimizer picked the correct PH of 2 for all these data sets.

As can be seen, for the four data set sizes, the optimal PH is 2. That is also the height calculated by our optimizer.

7.4 “Plate”

The third test case is another 2D stencil application called “Plate”. It is very similar to and replaces the “Poisson” application from the original test suite. It also models heat transfer in a plate but this time, the heat injection into the system is modeled as coming from the edge of the plate, not distributed throughout the plate. The makes the application different from Hotspot in two important ways. First, there is no required read of global data in each time step. Second, it gave us an opportunity to use the EdgeValue capability of the SL to inject the heat at the plate’s edge. Again the maximum square tile size on a GT200 card is 22x22 and valid PHs are 1 through 10. Again both the optimal PH and our model’s predicted PH is 2 for all four data set sizes. As the results for Plate are so similar to Hotspot, we don’t include a graph.

7.5 “Cell”

The fourth test case is a 3D stencil application called “Cell”. It models Conway’s Game of Life in 3 dimensions. In each time step, each cell calculates whether it is alive or dead based on the number of neighboring cells that were alive in the previous time step. To perform this calculation, it looks at its 26 nearest neighbors. For a 3D application, the maximum tile size on a GT200 card is 8x8x8, and the possible PHs are only 1 to 3. The runtimes vs. PH are shown in Figure 7. It is very clear from the chart that a PH of 3 is a very bad choice. The runtimes for PHs of 1 and 2 are similar, but 1 is better and is also predicted by our model.

7.6 Model Comparisons

As we have demonstrated, our model predicts the optimal PH for Hotspot, Plate, and Cell on the GTX-280 card. These were the same PHs predicted by the original model of Meng and Skadron for these same applications and data set sizes. However, they measured an optimal PH of 3 for Hotspot and Poisson. Their application implementations differ from ours in two important ways.

First, as their applications were hand optimized, they were able to load the read-only heat data used in Hotspot into shared memory to take advantage of temporal locality. Our SL has no way for the application programmer to express the access patterns for read-only

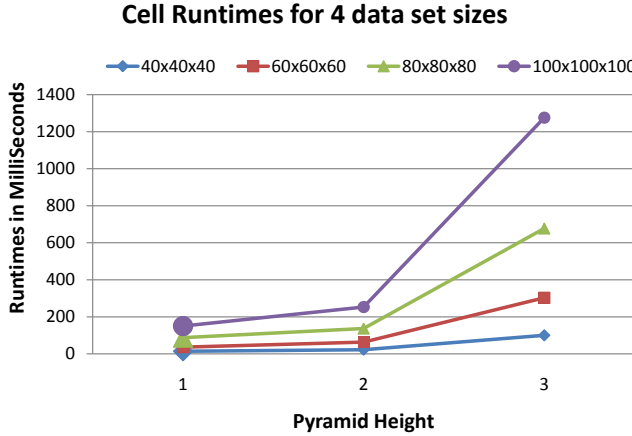


Figure 7. Runtimes for the Cell 3D stencil application vs. PH. The curve shows that a PH of 3 is a very bad choice. While the values for 1 and 2 are slightly compressed in the graph the optimal PH is in fact 1 for all data sets, as was accurately predicted by our optimizer.

data, and therefore our version of Hotspot will read the value from global memory in each time step. The ability of the hand optimized code to reuse the global memory read across several time steps tends to drive the optimal PH higher, leading to greater temporal reuse of data values loaded into shared memory. We hope to add better support for read-only data to SL at a future date.

Second, their inner loop code calculates the cell values for all cells in the tile, even those that have become too far away from the center of the tile to matter to the computation. Another option is to recalculate the zone of cells involved in useful computation on each time step. This keeps threads from calculating values that will never be used at the cost of more GZ size calculations. Given that our system must access read-only data from global memory in Hotspot (and potentially in other apps), the strategy of recalculating the GZ size in each time interval mitigates the number of global reads required, and is therefore the correct choice for such applications.

We found that 2D and 3D applications run faster by restricting the number of calculated values, even without using read only data. We believe this is because whole thread warps on the top and bottom of the tile will only have to check on each time step if they are still calculating useful values. In 3D apps, this effect is even more pronounced as thread warps in the front and back of the tile have this same advantage. The GZs on the sides of the tile tend not to span entire warps, so there is no advantage for warps at the sides of the tile from this effect. And the GZs in a 1D app only grow on the sides. Our CUDA templates take advantage of these effects by continuing to calculate values in the GZ for 1D stencil applications, but not for 2D or 3D applications. As a result of this change, our Pathfinder runtimes were significantly reduced.

We also analyzed the actual runtimes vs. our model predictions. Figure 8 shows Hotspot runtimes vs. predictions for PHs from 1 to 6. The model predictions are in terms of GPU cycles, not runtimes. Therefore, the model times are normalized to the actual runtime at the optimal PH=2. This allows a more direct comparison of the relative curves shapes, which is more important than absolute scaling for choosing the optimal PH. As the figure shows, the model has a steeper bowl shape than the actual runtimes. The effect is more pronounced for higher PHs, and really grow large quickly for PHs between 6 and 10.

The results presented here validate that the changes that we have made to the CUDA ISL analytical model have not lessened the




Figure 8. Runtimes for Hotspot vs. normalized runtimes predicted by the Model for PHs from 1 to 6. While the shape of the model curve does not exactly match the measured runtimes, the model accurately predicts an optimal PH of two.

effective accuracy of its PH predictions. Given the advantage of the elimination of the users’ manual involvement in this optimization, we believe this to be an important contribution of this work.

7.7 Comparison to Naïve CUDA code

In judging the performance of optimized code, it is important to see the effects of the optimizations compared against code that does not have those optimizations. Table 1 shows where we compare a “naïve” CUDA implementation for the four applications discussed above along with two additional applications.

The naïve CUDA implementation does not use shared memory. Instead, global device memory is always accessed. Therefore, each PE can only calculate one time step before synchronizing with neighboring PEs. While there are some edge cases to consider at the boundaries of the data set, there is no longer any need to handle edge cases across internal tile boundaries. However, the naïve code does still benefit from large tile sizes, and the maximum possible size is used. Overall, this code is much simpler than the optimized code, and would presumably be easier to write by hand.

Table 1 also shows the runtimes of the SL generated code, but with PH=1. Finally, the table also shows the runtime of the SL generated code with an “oracle” picking the best code. This often matches the SL optimized code, but differs in those cases where SL does not accurately predict the best PH, and the one case where the naïve code outperforms the SL generated code.

We can see from the table that the SL generated code outperforms the naïve code by between 2.37X and 0.97X with an arithmetic mean of 1.65X. The SL optimized code outperforms the naïve code by a significant factor for both the 1D Pathfinder application, and the 3D Cell application, but for different reasons. In Pathfinder, the optimal PH is high at 16, so there is substantial temporal reuse of the data values stored in shared memory even though the spatial reuse is limited by the single dimension of the data. This is further seen by looking at the PH=1 runtime for the SL code. By setting the PH=1, all of the temporal reuse is eliminated, and the overhead of the additional complexity of the SL code causes it to run slightly slower than the naïve code.

On the other hand, in Cell, the stencil involves 26 neighboring cell values. So, even though the temporal reuse is non-existent with an optimal PH of 1, the spatial reuse is significant. In the 2D applications, the optimal PH of Plate and Hotspot on the GTX-280

Application Data Set Size	Pathfinder 400K	Plate 1000 ²	PlateHalo 1000 ²	Plate++ 1000 ²	Hotspot 1000 ²	Cell 60 ³
Source of Best Time	O PH=16	SL PH=2	SL PH=2	SL PH=1	Naïve	SL PH=1
Oracle	9442 μ s	114172 μ s	131806 μ s	176292 μ s	140979 μ s	36505 μ s
SL Optimized vs Naïve	2.37	1.09	1.82	1.26	0.97	2.37
SL Optimized Time	9716 μ s	114172 μ s	131806 μ s	176292 μ s	145297 μ s	36505 μ s
SL PH=1	23591 μ s	145139 μ s	160101 μ s	176292 μ s	162902 μ s	36505 μ s
Naïve	23044 μ s	124930 μ s	239687 μ s	221964 μ s	140979 μ s	86600 μ s

Table 1. Runtimes for the SL optimized code for the four applications in Section 7, plus two new versions of Plate. Runtimes are also shown for a naïve CUDA implementation, along with a comparison relative to SL optimized code. For applications with large amounts of either temporal or spatial locality, the SL generated code far outperforms the naïve code. For other applications, the difference is more modest.

Application Data Set Size	Pathfinder 400K	Plate 1000 ²	PlateHalo 1000 ²	Plate++ 1000 ²	Hotspot 1000 ²	Cell 60 ³
Source of Best Time	O PH=32	SL PH=3	SL PH=3	SL PH=2	SL PH=2	SL PH=1
Oracle	8336 μ s	61367 μ s	66731 μ s	85854 μ s	77977 μ s	16044 μ s
SL Optimized vs Naïve	2.06	1.17	1.54	1.26	1.11	1.44
SL Optimized Time	8572 μ s	61367 μ s	66731 μ s	85854 μ s	77977 μ s	16044 μ s
SL PH=1	18373 μ s	81260 μ s	85148 μ s	91922 μ s	95763 μ s	16044 μ s
Naïve	17682 μ s	71844 μ s	102770 μ s	108047 μ s	86278 μ s	23048 μ s

Table 2. Data similar to Table 1 but now run on the Fermi series GTX-480 GPU card. Note that the increased block size on Fermi results in higher optimal PH for Plate, PlateHalo, and Plate++ due to the larger tile sizes. The inclusion of a cache on Fermi also results in a lower relative improvement of SL optimized code vs naïve for applications with large amounts of temporal or spatial locality.

are both 2, and the stencil involves 4 adjacent cells. So neither the spatial nor temporal reuse is extremely high. In fact the naïve code runs faster than SL’s optimal code for Hotspot, as the latter is also doing global memory reads in each time step.

To further explore these issues, we added two additional versions of Plate to our study. The first – called *PlateHalo* – expands the number of adjacent cells used in the computation to 8 by including the diagonal neighbors. This increases the spatial locality of the data in shared memory, and we see that the runtime for the naïve version almost doubles in PlateHalo over Plate, where the SL optimized version only slows by 15%. The SL optimized version of PlateHalo outperforms the naïve version by 1.8X.

The second variation of Plate is called *Plate++*. It also expands the stencil to include 8 neighbors, but this time by going out 2 cells to the north, south, east, and west. As this causes the GZs to grow twice as fast as in Plate, the optimal PH is now 1. So, temporal reuse is eliminated in the SL code, and SL code for Plate++ runs slower than either Plate or PlateHalo. The naïve code for Plate++ runs 8% faster than the naïve code for PlateHalo perhaps because of better memory access coalescence, or fewer memory bank conflicts.

7.8 Fermi Results

We reran the tests from Table 1 on a GTX-480, a Fermi series GPU. This machine contains a Core 2 Q9400 running at 2.66GHz with 6MB of L2 cache, and 5GB of DRAM. The system installation was otherwise identical to the GT200. On Fermi cards, the maximum threads per block is expanded to 1024, with 60 PEs each with 8 SIMD cores. The GTX-480 runs at 1.4GHz, and has a small global memory cache. The test results are shown in Table 2.

The increased numbers of threads per block means that the tile sizes for all applications are raised. As a result, the optimal PH for the applications tends to be higher because the GZ size for a given PH is a smaller percentage of the tile size. The optimal PH for Plate and PlateHalo have increased from 2 to 3, and for Plate++ is has increased from 1 to 2. For the shown data sizes, SL properly predicted these changes in optimal PH. However, for some of the other data set sizes SL did not. The worst prediction is for Cell with a data set size of 40*40*40 with a resultant slowdown of 83%. The

next worse was only 11%. The arithmetic mean of the slowdown over all tested applications and data set sizes was only 4.7%.

The inclusion of a cache on Fermi cards results in a reduction in the relative advantage of using the scratchpad memory for temporal and spatial locality. We see this in the reduced advantage of the SL optimized code over naïve in the 1D, 3D, and PlateHalo applications where these effects were largest on the GT200 architecture. Still, the SL optimized code performs better than the naïve code by between 2X and 1.1X with an arithmetic mean of 1.43X.

8. Related Work

Li et al. may have been the first to use the term *iterative stencil loops* [6]. They developed a compiler framework for these loops that used tiling to improve temporal data locality. Their work was intended to be run on a uniprocessor, so they did not face any communication costs and they did not consider ghost zones.

We are not the first to create a stencil compiler. Brickner et al. created a stencil compiler for the Connection Machine CM-2 massively parallel architecture [1]. The stencil expression was expressed in “microcode” which was specific to the CM-2.

In parallel with our efforts, Orchard et al. [14] have developed a stencil language called *Ypnos* as an extension to Haskell. They also argue for an architecture-independent declarative language for ISLs, and as in SL rely on language features instead of complex program analysis to determine the size and shape of the stencils and data. Their language includes reduction operators as a means for specifying convergence criteria, and they allow users to explicitly specify a PH for an ISL computation, but do not support automatic PH optimization. They do not report Ypnos performance results. Finally, our users specify the CellValue and EdgeValue calculations in C, as we feel it is more widely-accessible than Haskell.

Kamil et al. [4] studied the effects of tile size and time skewing to increase memory locality on the Itanium2, Opteron, Power5, and Cell. They concluded that controlling tile size, temporal locality across time slices, and hand-tuned control of scratchpad memory is critical to performance. We optimize use of CUDA’s scratchpad memory, tile size, and temporal reuse. However, our temporal lo-

cality is always strictly in the time dimension; we have not investigated use of time skewing. Datta et al. [3] studied optimizations for stencil applications across a wide range of architectures including CUDA. However, their optimizer uses auto-tuning to find optimal solutions. Our analytical model starting point allows our training period to be very short relative to a typical ISL computation.

Micikevicius explored hand coding stencil applications on CUDA, and ghost zone sizes [10]. They also explored using asynchronous data transfers for data sets too large to fit in GPU global memory. Volkov et al. explored hand optimizing CUDA programs involving parallel operations on large arrays [19]. We have included a subset of the optimizations presented by Micikevicius and Volkov in our CUDA template library, but we have no support for arrays too large to fit into the devices global memory, and therefore we did not investigate asynchronous data transfers.

Several attempts have been made to build simplifying layers on top of the CUDA language. Lee et al. built a source-to-source translator for OpenMP to CUDA [5]. OpenMP can express a significantly wider range of parallel computations than that which can be expressed by SL. However, this also means an OpenMP translator cannot include stencil application specific optimizations as we do.

CUDA-lite is a general purpose source-to-source translator for automatically optimizing CUDA programs [18]. The user specifies a naïve CUDA implementation of their ISL that only uses global memory, and adds some pragmas that give hints about the size, shape, and data access patterns. Then CUDA-lite generates an optimized CUDA program that leverages the specialized memory spaces, resulting in 2-17X speedups for their test applications. It appears they ran their tests on earlier GPUs that had much less flexible memory access coalescing hardware than modern GPUs. In addition, their primary test applications were not stencil applications, and instead performed operations such as matrix multiply. These contain large amounts of spatial and temporal locality, and are compute intensive, allowing for more processing time to hide memory access latency [15]. By comparison, SL also optimizes for shared memory but not texture memory, and also optimizes PH.

9. Conclusions and Future Directions

We have defined a new language, SL, for easily expressing stencil applications. SL is a formal high-level language in which a programmer can conveniently describe an ISL computation by focusing on the simple computation needed to compute cell values. No optimization nor architectural-specific information is needed.

We have created a SL compiler that uses templates to allow the generated code to be extended to many parallel execution environments. Based on a single source file, it could output efficient code for CUDA, OpenCL, OpenMP, pthreads, or other environments.

We have implemented an initial template and an optimizer for the CUDA programming environment that blocks computations to use the scratchpad memory, and that estimates the best PH. The optimizer is fully automatic and utilizes both static and dynamic application information. It accurately calculates the optimal tile size and PHs for several test applications, on both GT200 and Fermi series GPUs. As a result, with an SL declaration of about 16 lines, a user can achieve average speedups of 1.65X% for the GTX-280, and 1.43X for the GTX-480, over naïve CUDA code. Yet the SL specification is much easier to write than even the naïve CUDA code, and is similar to writing code for a single thread CPU.

The optimizer's ability to leverage dynamically acquired information about the runtime environment means that the generated code can be used on a variety of devices. Our model ported directly from the GT200 architecture to the next generation Fermi architecture with only trivial modification.

The current SL architecture can be extended in several directions. First, with language extensions, we can improve support for

read-only data at least in the common cases such as Hotspot where the array access pattern is straightforward. A second SL language extension could add support for reduction operators, and ISL termination based on a data convergence criteria. It would also be interesting to investigate data morphing as a way to improve cache performance or reduce memory bank conflicts.

References

- [1] R. Brickner, W. George, S. L. Johnsson, and A. Ruttenberg. A Stencil Compiler for the Connection Machine Models CM-2/200. In *Fourth Workshop on Compilers for Parallel Computers*, 1993.
- [2] B. Chapman, G. Jost, and R. v. d. Pas. *Using OpenMP: Portable Shared Memory Parallel Programming (Scientific and Engineering Computation)*. The MIT Press, 2007.
- [3] K. Datta, M. Murphy, V. Volkov, S. Williams, J. Carter, L. Oliker, D. Patterson, J. Shalf, and K. Yelick. Stencil Computation Optimization and Auto-tuning on State-of-the-art Multicore Architectures. In *Conference on Supercomputing*, pages 1–12, Piscataway, NJ, 2008.
- [4] S. Kamil, K. Datta, S. Williams, L. Oliker, J. Shalf, and K. Yelick. Implicit and Explicit Optimizations for Stencil Computations. In *Workshop on Memory System Performance and Correctness*, pages 51–60, 2006.
- [5] S. Lee, S.-J. Min, and R. Eigenmann. OpenMP to GPGPU: A Compiler Framework for Automatic Translation and Optimization. *SIGPLAN Notices*, 44(4):101–110, 2009.
- [6] Z. Li and Y. Song. Automatic Tiling of Iterative Stencil Loops. *ACM TOPLAS*, 26(6):975–1028, 2004.
- [7] E. Lindholm, J. Nickolls, S. Oberman, and J. Montrym. NVIDIA Tesla: A Unified Graphics and Computing Architecture. *IEEE Micro*, 28(2):39–55, 2008.
- [8] J. R. Mashey. War of the Benchmark Means: Time for a Truce. *SIGARCH Computer Architecture News*, 32(4):1–14, 2004.
- [9] J. Meng and K. Skadron. Performance Modeling and Automatic Ghost Zone Optimization for Iterative Stencil Loops on GPUs. In *23rd International Conference on Supercomputing*, pages 256–265, 2009.
- [10] P. Micikevicius. 3D Finite Difference Computation on GPUs using CUDA. In *2nd Workshop on General Purpose Processing on Graphics Processing Units*, pages 79–84, 2009.
- [11] B. Nichols, D. Buttlar, and J. P. Farrell. *Pthreads Programming*. O'Reilly & Associates, Inc., Sebastopol, CA, 1996.
- [12] J. Nickolls and W. Dally. The GPU Computing Era. *IEEE Micro*, 30(2):56–69, March-April 2010.
- [13] J. Nickolls, I. Buck, M. Garland, and K. Skadron. Scalable Parallel Programming with CUDA. *Queue*, 6(2):40–53, 2008.
- [14] D. A. Orchard, M. Bolingbroke, and A. Mycroft. Ypnos: Declarative, Parallel Structured Grid Programming. In *5th Workshop on Declarative Aspects of Multicore Programming*, pages 15–24, New York, 2010.
- [15] S. Ryoo, C. I. Rodrigues, S. S. Baghsorkhi, S. S. Stone, D. B. Kirk, and W.-m. W. Hwu. Optimization Principles and Application Performance Evaluation of a Multithreaded GPU using CUDA. In *Principles and Practice of Parallel Programming*, pages 73–82, New York, 2008.
- [16] K. Skadron, M. R. Stan, K. Sankaranarayanan, W. Huang, S. Velusamy, and D. Tarjan. Temperature-Aware Microarchitecture: Modeling and Implementation. *ACM TACO*, 1(1):94–125, 2004.
- [17] J. Stone, D. Gohara, and G. Shi. OpenCL: A Parallel Programming Standard for Heterogeneous Computing Systems. *Computing in Science Engineering*, 12(3):66–73, May-June 2010.
- [18] S.-Z. Ueng, M. Lathara, S. S. Baghsorkhi, and W.-M. W. Hwu. CUDA-Lite: Reducing GPU Programming Complexity. In *21st Workshop on Languages and Compilers for Parallel Computing*, pages 1–15, Berlin, Heidelberg, 2008.
- [19] Volkov, Vasily and Demmel, James. LU, QR and Cholesky Factorizations using Vector Capabilities of GPUs. Technical Report UCB/EECS-2008-49, University of California at Berkeley, May 2008.